

OPENSEARCH LOGGING

INTERN PROJECT BY EKENE ILLOH

BACKGROUND

- ADAM-Galaxy has several services (Arches, AIS, GRS, GAMS etc.) involved in registering an advertising account
- Sometimes, tracking the status of an account registration can be very difficult since the logs for the services are stored in different places, formats and systems
- It becomes difficult and time-consuming to query the different databases but all of them frequently have critical information about an account registration
- The goal was to build an OpenSearch cluster that will aggregate application logs from account registration from adam services into an easily queryable place

HOW THIS WOULD WORK – ELK STACK

- The logs are transferred from the services to Elasticsearch through the use of the ELK stack – Elasticsearch, Logstash and Kibana
- Using Logstash you can parse and push the entries in log files to Elasticsearch
- Kibana gives you a powerful web interface with search, filter, and statistical analysis functionality driven by Elasticsearch
- Using this stack we can have real-time access to all log entries from all hosts running a service

RiverPig6



LogStash



LogStash.conf tells LogStash where to find the logs and what to do with them

LogStash.conf

LogStashProps.json

Elasticsearch



Kibana



Cognito

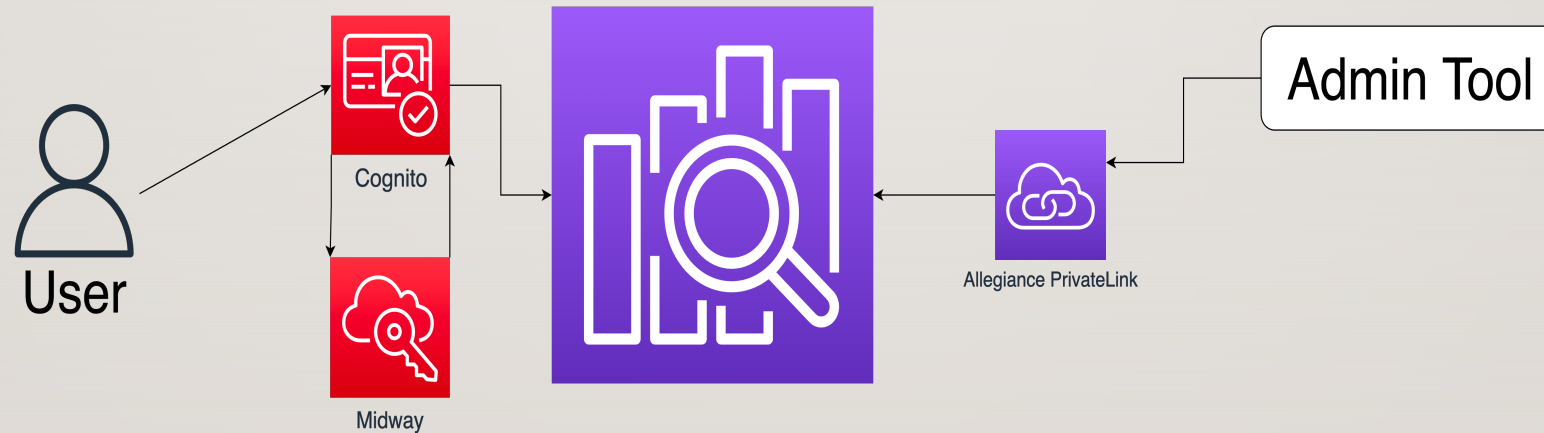


Controls access to Kibana

Amazon Federate

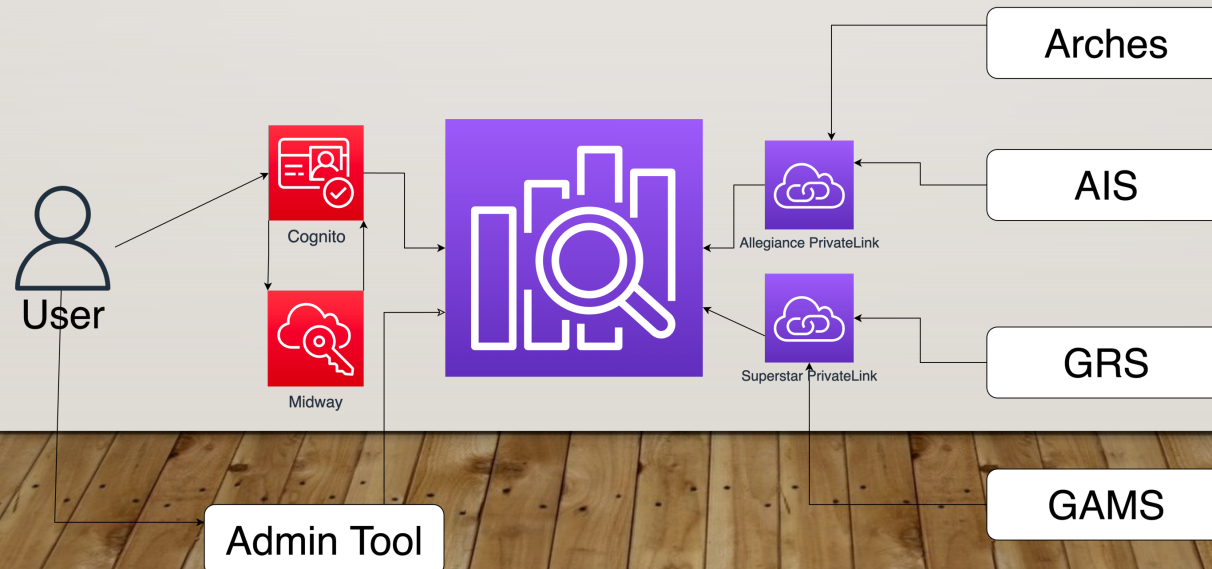
EXISTING SYSTEM

- The POC had Admin Tool onboarded to the cluster through the allegiance private link
- The cluster was set up to receive logs from Admin Tool beta NA



OBJECTIVE

- Expand the current POC to take in logs from the following additional services: GRS, GAMS, AIS and Arches
 - Scaled down to just AIS and Arches
- Build a programmatic API that would allow a user to query a service in the opensearch cluster to retrieve all the logs for a given registration attempt



PROCESS

- Set up the existing POC
- Update the Admin Tool package to onboarding beta logs from all regions
- Make changes to ADAMTimberCDK and GlobalAccountsConstruct packages to allow the onboarding of additional services
- Start onboarding the additional services by making the code changes to their packages

DEMO



CHALLENGES

- Unfamiliar tech stack
- Backend development
- Debugging build errors
- Deploying the ADAM Timber CDK

NEXT STEPS

- Clean up CDK codebase
- Set up endpoints in vpc endpoint service using CDK instead of manually on AWS console
- Build an API for querying the logs in OpenSearch
- Expand the cluster to receive logs from GRS and GAMS
- Expand the system from beta to prod for all the services

THANK YOU!